

Responsible AI Impact Assessment for JobAssistAI

For questions about specific sections within the Impact Assessment, please refer to the Impact Assessment Guide.

Section 1: System Information

System profile

1.1 Complete the system information below.

System name	JobAssistAI
Team name	JobAssistAI

Track revision history below.

Authors	Joshua Kaelin
Last updated	2025-03-19

Identify the individuals who will review your Impact Assessment when it is completed.

Reviewers	Harshit Jain Srujana Vanama
-----------	--------------------------------

System lifecycle stage

1.2 Indicate the dates of planned releases for the system.

Date	Lifecycle stage
2025-03-11	Planning & analysis
2025-03-11	Design
2025-03-11	Development
2025-03-15	Testing

2025-03-17	Implementation & deployment
Unknown	Maintenance
Unknown	Retired

System description

1.3 Briefly explain, in plain language, what you're building. This will give reviewers the necessary context to understand the system and the environment in which it operates.

System description

We are building a system to assist job coaches in their mission of supporting disabled persons find employment. The initial release will be for demo purposes only and include only mock data with no connection to real personal information. If implemented in a production environment, the system would need to integrate with the workflow management system of the supported agency, generally managed by local government.

If you have links to any supplementary information on the system such as demonstrations, functional specifications, slide decks, or system architecture diagrams, please include links below.

Description of supplementary information	Link
GitHub repository	https://github.com/harshitjain17/JobAssist-AI

System purpose

1.4 Briefly describe the purpose of the system and system features, focusing on how the system will address the needs of the people who use it. Explain how the AI technology contributes to achieving these objectives.

System purpose

The purpose of the system is to assist job coaches in their mission mission of supporting disabled persons find employment.

System features

1.5 Focusing on the whole system, briefly describe the system features or high-level feature areas that already exist and those planned for the upcoming release.

Existing system features	System features planned for the upcoming release
	AI Assistant (chat + events)
	AI Document processing
	AI Knowledge Base

Briefly describe how this system relates to other systems or products. For example, describe if the system includes models from other systems.

Relation to other systems/products
None at this time.

Geographic areas and languages

1.6 Describe the geographic areas where the system will or might be deployed to identify special considerations for language, laws, and culture.

The system is currently deployed to:	Azure east-us & east-us2
In the upcoming release, the system will be deployed to:	Azure east-us & east-us2
In the future, the system might be deployed to:	Any Azure US region

For natural language processing systems, describe supported languages:

The system currently supports:	English
In the upcoming release, the system will support:	English
In the future, the system might support:	The AI features will generally be language-agnostic per their underlying model support.

Deployment mode

1.7 Document each way that this system might be deployed.

How is the system currently deployed?	GitHub Actions
---------------------------------------	----------------

Will the deployment mode change in the upcoming release?
If so, how?

GitHub Actions

Intended uses

1.8 *Intended uses are the uses of the system your team is designing and testing for. An intended use is a description of who will use the system, for what task or purpose, and where they are when using the system. They are not the same as system features, as any number of features could be part of an intended use. Fill in the table with a description of the system's intended use(s).*

Name of intended use(s)	Description of intended use(s)
1. AI Assistant	Provide general information about the user's roles & responsibilities, the official procedures to which they must adhere and guidance based on the user's current context.
2. Document Processing	Converting handwritten notes to structured, context-aware objects & generating PDF documents with the structured data.
3. Knowledge Base	Allow users to share free-form knowledge via text or voice memos.

Section 2A: Intended uses - AI Assistant

Intended use #1: AI Assistant

Assessment of fitness for purpose

2.1 Assess how the system's use will solve the problem posed by each intended use, recognizing that there may be multiple valid ways in which to solve the problem.

Assessment of fitness for purpose

A job coach's role is broad and support is limited. They are overburdened with bureaucratic, administrative tasks. The AI Assistant will seek to create efficiencies in their workflows by giving them specific, context-aware advice on actions they may want to take or answer general questions based on their department's documents.

As a broad disclaimer regarding the fitness of this system, this initial release was developed during a hackathon and certain due diligence tasks regarding fitness (user interviews, workflow assessments, etc.) were not undertaken but the JobAssist AI team recognizes these as critical-path items to any successful deployment.

Stakeholders, potential benefits, and potential harms

2.2 Identify the system's stakeholders for this intended use. Then, for each stakeholder, document the potential benefits and potential harms. For more information, including prompts, see the *Impact Assessment Guide*.

Stakeholders	Potential system benefits	Potential system harms
1. Job Coaches	More efficient workflow, less need to track information outside the system.	Already-overwhelmed job coaches will have something new to learn. May be seen as decreased autonomy if more of their work is reflected in the system.
2. Job Consumers	Improved hand-offs between coaches, more consistent processes.	Poor performance of the AI Assistant could potentially lead to suboptimal outcomes.
3. Local Government	Potential cost-avoidance based on efficiency improvements. KPI improvements.	There is a change management cost associated with any change.
4. Participating Employers	Job Coach efficiency should directly translate into better support & more rapidly filling open positions.	

Stakeholders for Goal-driven requirements from the Responsible AI Standard

2.3 Certain Goals in the Responsible AI Standard require you to identify specific types of stakeholders. You may have included them in the stakeholder table above. For the Goals below that apply to the system, identify the specific stakeholder(s) for this intended use. If a Goal does not apply to the system, enter "N/A" in the table.

Goal A5: Human oversight and control

This Goal applies to all AI systems. Complete the table below.

Who is responsible for troubleshooting, managing, operating, overseeing, and controlling the system during and after deployment?	For these stakeholders, identify their oversight and control responsibilities.
Initially, these functions would need to be managed by the JobAssistAI team. Ideally, certain aspects of the system would become more configurable by local admins within the target organization.	<p>The primary responsibilities are:</p> <ul style="list-style-type: none"> Refining the system prompt to better reflect the environment in which it is being used. Determining the source documents to be indexed and the update mechanism for these sources. Monitoring usage of the system for abuse, key metrics & user satisfaction.

Goal T1: System intelligibility for decision making

This Goal applies to AI systems when the intended use of the generated outputs is to inform decision making by or about people. If this Goal applies to the system, complete the table below.

Who will use the outputs of the system to make decisions?	Who will decisions be made about?
Job Coaches	<p>Primarily: job consumers.</p> <p>Secondarily: employers, potential employers, adjacent agencies, community centers, etc.</p>

Goal T2: Communication to stakeholders

This Goal applies to all AI systems. Complete the table below.

Who will make decisions about whether to employ the system for particular tasks?	Who develops or deploys systems that integrate with this system?
Ultimately, the Job Coach will make the decision whether or not to use the system. However, it cannot	The JobAssistAI team would play a role in initial integrations while a development/integration framework is created. Ultimately, anyone assigned

be ruled out that use of the system may become mandatory per organizational policy.

to develop or deploy related systems on behalf of these agencies would be included here.

Goal T3: Disclosure of AI interaction

This Goal applies to AI systems that impersonate interactions with humans, unless it is obvious from the circumstances or context of use that an AI system is in use, and AI systems that generate or manipulate image, audio, or video content that could falsely appear to be authentic. If this Goal applies to the system, complete the table below.

Who will use or be exposed to the system?

Per the current scope, only trained Job Coaches will directly use the system.
All AI interactions are accompanied by associated "bot" iconography and "AI" text headings.

Fairness considerations

2.4 For each Fairness Goal that applies to the system, 1) identify the relevant stakeholder(s) (e.g., system user, person impacted by the system); 2) identify any demographic groups, including marginalized groups, that may require fairness considerations; and 3) prioritize these groups for fairness consideration and explain how the fairness consideration applies. If the Fairness Goal does not apply to the system, enter "N/A" in the first column.

Goal F1: Quality of service

This Goal applies to AI systems when system users or people impacted by the system with different demographic characteristics might experience differences in quality of service that can be remedied by building the system differently. If this Goal applies to the system, complete the table below describing the appropriate stakeholders for this intended use.

Which stakeholder(s) will be affected?	For affected stakeholder(s) which demographic groups are you prioritizing for this Goal?	Explain how each demographic group might be affected.
Job Consumers	Disabled Job Consumers, i.e. job consumers with one or more impairments relating to their vision, mobility, mental status, etc.	The quality of service for Disabled Job Consumers is largely based on the performance of their Job Coach. By improving the efficiency & policy adherence of the Job Coach, we seek to improve the consumer's quality of service.

Goal F2: Allocation of resources and opportunities

This Goal applies to AI systems that generate outputs that directly affect the allocation of resources or opportunities relating to finance, education, employment, healthcare, housing, insurance, or social welfare. If this Goal applies to the system, complete the table below describing the appropriate stakeholders for this intended use.

Which stakeholder(s) will be affected?	For affected stakeholder(s) which demographic groups are you prioritizing for this Goal?	Explain how each demographic group might be affected.
Job Consumers	Disabled Job Consumers, i.e. job consumers with one or more impairments relating to their vision, mobility, mental status, etc.	Impacting employment decisions is critical path functionality for this system. We seek to improve the process of getting job consumers from their intake through successful job placement.

Goal F3: Minimization of stereotyping, demeaning, and erasing outputs

This Goal applies to AI systems when system outputs include descriptions, depictions, or other representations of people, cultures, or society. If this Goal applies to the system, complete the table below describing the appropriate stakeholders for this intended use.

Which stakeholder(s) will be affected?	For affected stakeholder(s) which demographic groups are you prioritizing for this Goal?	Explain how each demographic group might be affected.
Job Consumers	Disabled Job Consumers, i.e. job consumers with one or more impairments relating to their vision, mobility, mental status, etc.	We consider this stakeholder group to be a vulnerable population and, accordingly, we need take precautions to minimize any stereotyping of persons with disabilities, instead focusing on the abilities of the individual job consumer.

Technology readiness assessment

2.5 Indicate with an "X" the description that best represents the system regarding this intended use.

Select one	Technology Readiness
	The system includes AI supported by basic research and has not yet been deployed to production systems at scale for similar uses.

X	The system includes AI supported by evidence demonstrating feasibility for uses similar to this intended use in production systems.
	This is the first time that one or more system component(s) are to be validated in relevant environment(s) for the intended use. Operational conditions that can be supported have not yet been completely defined and evaluated.
	This is the first time the whole system will be validated in relevant environment(s) for the intended use. Operational conditions that can be supported will also be validated. Alternatively, nearly similar systems or nearly similar methods have been applied by other organizations with defined success.
	The whole system has been deployed for all intended uses , and operational conditions have been qualified through testing and uses in production.

Task complexity

2.6 Indicate with an "X" the description that best represents the system regarding this intended use.

Select One	Task Complexity
	Simple tasks , such as classification based on few features into a few categories with clear boundaries. For such decisions, humans could easily agree on the correct answer, and identify mistakes made by the system. For example, a natural language processing system that checks spelling in documents.
X	Moderately complex tasks , such as classification into a few categories that are subjective. Typically, ground truth is defined by most evaluators arriving at the same answer. For example, a natural language processing system that autocompletes a word or phrase as the user is typing.
	Complex tasks , such as models based on many features, not easily interpretable by humans, resulting in highly variable predictions without clear boundaries between decision criteria. For such decisions, humans would have a difficult time agreeing on the best answer, and there may be no clearly incorrect answer. For example, a natural language processing system that generates prose based on user input prompts.

Role of humans

2.7 Indicate with an "X" the description that best represents the system regarding this intended use.

Select One	Role of humans
------------	----------------

	People will be responsible for troubleshooting triggered by system alerts but will not otherwise oversee system operation. For example, an AI system that generates keywords from unstructured text alerts the operator of errors, such as improper format of submission files.
	The system will support effective hand-off to people but will be designed to automate most use. For example, an AI system that generates keywords from unstructured text that can be configured by system admins to alert the operator when keyword generation falls below a certain confidence threshold.
	The system will require effective hand-off to people but will be designed to automate most use. For example, an AI system that generates keywords from unstructured text alerts the operator when keyword generation falls below a certain confidence threshold (regardless of system admin configuration).
	People will evaluate system outputs and can intervene before any action is taken: the system will proceed unless the reviewer intervenes. For example, an AI system that generates keywords from unstructured text will deliver the generated keywords for operator review but will finalize the results unless the operator intervenes.
X	People will make decisions based on output provided by the system: the system will not proceed unless a person approves. For example, an AI system that generates keywords from unstructured text but does not finalize the results without review and approval from the operator.

Deployment environment complexity

2.8 Indicate with an "X" the description that best represents the system regarding this intended use.

Select One	Deployment environment complexity
	Simple environment , such as when the deployment environment is static, possible input options are limited, and there are few unexpected situations that the system must deal with gracefully. For example, a natural language processing system used in a controlled research environment.
X	Moderately complex environment , such as when the deployment environment varies, unexpected situations the system must deal with gracefully may occur, but when they do, there is little risk to people, and it is clear how to effectively mitigate issues. For example, a natural language processing system used in a corporate workplace where language is professional and communication norms change slowly.

	<p>Complex environment, such as when the deployment environment is dynamic, the system will be deployed in an open and unpredictable environment or may be subject to drifts in input distributions over time. There are many possible types of inputs, and inputs may significantly vary in quality. Time and attention may be at a premium in making decisions and it can be difficult to mitigate issues. For example, a natural language processing system used on a social media platform where language and communication norms change rapidly.</p>
--	--

Section 2B: Intended uses - Document Processing

Intended use #2: Document Processing

Assessment of fitness for purpose

2.1 Assess how the system's use will solve the problem posed by each intended use, recognizing that there may be multiple valid ways in which to solve the problem.

Assessment of fitness for purpose

A job coach's role is broad and support is limited. They are overburdened with bureaucratic, administrative tasks.

The Document Processing capability will seek to create efficiencies in their workflows by:

- Allowing users to use photos (e.g. JPEG file from phone camera) of handwritten notes
- Convert handwritten notes into structured text & context
- Produce output documents (e.g. PDF files formatted per organizational policy)

As a broad disclaimer regarding the fitness of this system, this initial release was developed during a hackathon and certain due diligence tasks regarding fitness (user interviews, workflow assessments, etc.) were not undertaken but the JobAssist AI team recognizes these as critical-path items to any successful deployment.

Stakeholders, potential benefits, and potential harms

2.2 Identify the system's stakeholders for this intended use. Then, for each stakeholder, document the potential benefits and potential harms. For more information, including prompts, see the Impact Assessment Guide.

Stakeholders	Potential system benefits	Potential system harms
1. Job Coaches	More efficient workflow, less need to track information outside the system.	Poor performance of the system could discourage its use and sabotage future efforts to streamline these workflows.
2. Job Consumers	Job Consumers should see benefits based on having more data related to their case within the system for tracking & decision-making.	Job Coaches with substantially different usage of this feature may lead to relative gaps in service for their respective Consumers.
3. Local Government	Standardization of workflows is a foundational capability to meeting policy adherence goals.	There is a change management cost associated with any change.

Stakeholders for Goal-driven requirements from the Responsible AI Standard

2.3 Certain Goals in the Responsible AI Standard require you to identify specific types of stakeholders. You may have included them in the stakeholder table above. For the Goals below that apply to the system, identify the specific stakeholder(s) for this intended use. If a Goal does not apply to the system, enter "N/A" in the table.

Goal A5: Human oversight and control

This Goal applies to all AI systems. Complete the table below.

Who is responsible for troubleshooting, managing, operating, overseeing, and controlling the system during and after deployment?	For these stakeholders, identify their oversight and control responsibilities.
Initially, these functions would need to be managed by the JobAssistAI team. Ideally, certain aspects of the system would become more configurable by local admins within the target organization.	<p>The primary responsibilities are:</p> <ul style="list-style-type: none"> Establishing the primary workflow and any required additional workflows. Determining supported document sources. Configuring parameters for document processing & context inference. Configuring output document templates & available capabilities. Determining the privacy level and/or sharing policies available. Monitoring usage of the system for abuse, key metrics & user satisfaction.

Goal T1: System intelligibility for decision making

This Goal applies to AI systems when the intended use of the generated outputs is to inform decision making by or about people. If this Goal applies to the system, complete the table below.

Who will use the outputs of the system to make decisions?	Who will decisions be made about?
<p>Directly: none</p> <p>Indirectly: Job Consumers</p>	<p>This is generally a process efficiency improvement capability and the outputs are to meet administrative needs, not directly inform decisions.</p> <p>However, future capabilities may associate this document data with a job consumer and index it in a manner that would allow it to be incorporated into decision-making processes such as 2A.</p>

Goal T2: Communication to stakeholders

This Goal applies to all AI systems. Complete the table below.

Who will make decisions about whether to employ the system for particular tasks?	Who develops or deploys systems that integrate with this system?
Ultimately, the Job Coach will make the decision whether or not to use the system. However, it cannot be ruled out that use of the system may become mandatory per organizational policy.	The JobAssistAI team would play a role in initial integrations while a development/integration framework is created. Ultimately, anyone assigned to develop or deploy related systems on behalf of these agencies would be included here.

Goal T3: Disclosure of AI interaction

This Goal applies to AI systems that impersonate interactions with humans, unless it is obvious from the circumstances or context of use that an AI system is in use, and AI systems that generate or manipulate image, audio, or video content that could falsely appear to be authentic. If this Goal applies to the system, complete the table below.

Who will use or be exposed to the system?
As a process improvement / automation, there is limited view into the process from any system user. From the Job Coach's perspective, this is simply an automation.

Fairness considerations

2.4 *For each Fairness Goal that applies to the system, 1) identify the relevant stakeholder(s) (e.g., system user, person impacted by the system); 2) identify any demographic groups, including marginalized groups, that may require fairness considerations; and 3) prioritize these groups for fairness consideration and explain how the fairness consideration applies. If the Fairness Goal does not apply to the system, enter "N/A" in the first column.*

Goal F1: Quality of service

This Goal applies to AI systems when system users or people impacted by the system with different demographic characteristics might experience differences in quality of service that can be remedied by building the system differently. If this Goal applies to the system, complete the table below describing the appropriate stakeholders for this intended use.

Which stakeholder(s) will be affected?	For affected stakeholder(s) which demographic groups are you prioritizing for this Goal?	Explain how each demographic group might be affected.

Job Coaches	For the purposes of the initial demo release, our presumptive Job Coach is based on the information we were provided for the hackathon. We are fully aware that each Job Coach is an individual first and we must always seek to meet the user where they are.	The quality of service for Job Coaches is critical to the adoption of this feature. Over-optimizing the system for a hypothetical Job Coach will likely under-serve individual users. It will be important to revisit our assumptions based on user interviews.
-------------	--	---

Goal F2: Allocation of resources and opportunities

This Goal applies to AI systems that generate outputs that directly affect the allocation of resources or opportunities relating to finance, education, employment, healthcare, housing, insurance, or social welfare. If this Goal applies to the system, complete the table below describing the appropriate stakeholders for this intended use.

Which stakeholder(s) will be affected?	For affected stakeholder(s) which demographic groups are you prioritizing for this Goal?	Explain how each demographic group might be affected.
N/A		

Goal F3: Minimization of stereotyping, demeaning, and erasing outputs

This Goal applies to AI systems when system outputs include descriptions, depictions, or other representations of people, cultures, or society. If this Goal applies to the system, complete the table below describing the appropriate stakeholders for this intended use.

Which stakeholder(s) will be affected?	For affected stakeholder(s) which demographic groups are you prioritizing for this Goal?	Explain how each demographic group might be affected.
Job Consumers	Disabled Job Consumers, i.e. job consumers with one or more impairments relating to their vision, mobility, mental status, etc.	We consider this stakeholder group to be a vulnerable population and, accordingly, we need take precautions to minimize any stereotyping of persons with disabilities, instead focusing on the abilities of the individual job consumer. To that end, we need to consider what this process do if it encounters potentially improper language or comments in the source notes.

Technology readiness assessment

2.5 Indicate with an "X" the description that best represents the system regarding this intended use.

Select one	Technology Readiness
	The system includes AI supported by basic research and has not yet been deployed to production systems at scale for similar uses.
X	The system includes AI supported by evidence demonstrating feasibility for uses similar to this intended use in production systems.
	This is the first time that one or more system component(s) are to be validated in relevant environment(s) for the intended use. Operational conditions that can be supported have not yet been completely defined and evaluated.
	This is the first time the whole system will be validated in relevant environment(s) for the intended use. Operational conditions that can be supported will also be validated. Alternatively, nearly similar systems or nearly similar methods have been applied by other organizations with defined success.
	The whole system has been deployed for all intended uses , and operational conditions have been qualified through testing and uses in production.

Task complexity

2.6 Indicate with an "X" the description that best represents the system regarding this intended use.

Select One	Task Complexity
	Simple tasks , such as classification based on few features into a few categories with clear boundaries. For such decisions, humans could easily agree on the correct answer, and identify mistakes made by the system. For example, a natural language processing system that checks spelling in documents.
X	Moderately complex tasks , such as classification into a few categories that are subjective. Typically, ground truth is defined by most evaluators arriving at the same answer. For example, a natural language processing system that autocompletes a word or phrase as the user is typing.

	<p>Complex tasks, such as models based on many features, not easily interpretable by humans, resulting in highly variable predictions without clear boundaries between decision criteria. For such decisions, humans would have a difficult time agreeing on the best answer, and there may be no clearly incorrect answer. For example, a natural language processing system that generates prose based on user input prompts.</p>
--	--

Role of humans

2.7 Indicate with an "X" the description that best represents the system regarding this intended use.

Select One	Role of humans
	<p>People will be responsible for troubleshooting triggered by system alerts but will not otherwise oversee system operation. For example, an AI system that generates keywords from unstructured text alerts the operator of errors, such as improper format of submission files.</p>
	<p>The system will support effective hand-off to people but will be designed to automate most use. For example, an AI system that generates keywords from unstructured text that can be configured by system admins to alert the operator when keyword generation falls below a certain confidence threshold.</p>
X	<p>The system will require effective hand-off to people but will be designed to automate most use. For example, an AI system that generates keywords from unstructured text alerts the operator when keyword generation falls below a certain confidence threshold (regardless of system admin configuration).</p>
	<p>People will evaluate system outputs and can intervene before any action is taken: the system will proceed unless the reviewer intervenes. For example, an AI system that generates keywords from unstructured text will deliver the generated keywords for operator review but will finalize the results unless the operator intervenes.</p>
	<p>People will make decisions based on output provided by the system: the system will not proceed unless a person approves. For example, an AI system that generates keywords from unstructured text but does not finalize the results without review and approval from the operator.</p>

Deployment environment complexity

2.8 Indicate with an "X" the description that best represents the system regarding this intended use.

Select One	Deployment environment complexity
------------	-----------------------------------

	<p>Simple environment, such as when the deployment environment is static, possible input options are limited, and there are few unexpected situations that the system must deal with gracefully. For example, a natural language processing system used in a controlled research environment.</p>
X	<p>Moderately complex environment, such as when the deployment environment varies, unexpected situations the system must deal with gracefully may occur, but when they do, there is little risk to people, and it is clear how to effectively mitigate issues. For example, a natural language processing system used in a corporate workplace where language is professional and communication norms change slowly.</p>
	<p>Complex environment, such as when the deployment environment is dynamic, the system will be deployed in an open and unpredictable environment or may be subject to drifts in input distributions over time. There are many possible types of inputs, and inputs may significantly vary in quality. Time and attention may be at a premium in making decisions and it can be difficult to mitigate issues. For example, a natural language processing system used on a social media platform where language and communication norms change rapidly.</p>

Section 2C: Intended uses – Knowledge Base

Intended use #3: Knowledge Base

Assessment of fitness for purpose

2.1 Assess how the system's use will solve the problem posed by each intended use, recognizing that there may be multiple valid ways in which to solve the problem.

Assessment of fitness for purpose

A job coach's role is broad and support is limited. They are overburdened with bureaucratic, administrative tasks.

The Knowledge Base capability will seek to create a sustainable, searchable & durable repository of information, the bulk of which is not stored in any standardized way in the current system.

Primary Features:

- Allow Job Coaches to submit unstructured text snippets or voice memos to the Knowledge Base
- Allow Job Coaches to query the Knowledge Base using natural language or keywords.
- Future: Allow [Intended use #1: AI Assistant] to incorporate the Knowledge Base as a data source

As a broad disclaimer regarding the fitness of this system, this initial release was developed during a hackathon and certain due diligence tasks regarding fitness (user interviews, workflow assessments, etc.) were not undertaken but the JobAssist AI team recognizes these as critical-path items to any successful deployment.

Stakeholders, potential benefits, and potential harms

2.2 Identify the system's stakeholders for this intended use. Then, for each stakeholder, document the potential benefits and potential harms. For more information, including prompts, see the Impact Assessment Guide.

Stakeholders	Potential system benefits	Potential system harms
1. Job Coaches	More efficient workflow, easier to locate information distributed in multiple repositories (digital, written, informally "known")	Poor performance of the system could discourage its use and sabotage future efforts to streamline these workflows.
2. Job Consumers	Improved hand-offs between coaches, more consistent processes.	Job Coaches with substantially different usage of this feature may lead to relative gaps in service for their respective Consumers.
3. Local Government	Standardization of workflows is a foundational capability to meeting policy adherence goals. Improve ease & consistency of	There is a change management cost associated with any change.

	transitioning consumers between coaches.	
--	--	--

Stakeholders for Goal-driven requirements from the Responsible AI Standard

2.3 Certain Goals in the Responsible AI Standard require you to identify specific types of stakeholders. You may have included them in the stakeholder table above. For the Goals below that apply to the system, identify the specific stakeholder(s) for this intended use. If a Goal does not apply to the system, enter "N/A" in the table.

Goal A5: Human oversight and control

This Goal applies to all AI systems. Complete the table below.

Who is responsible for troubleshooting, managing, operating, overseeing, and controlling the system during and after deployment?	For these stakeholders, identify their oversight and control responsibilities.
Initially, these functions would need to be managed by the JobAssistAI team. Ideally, certain aspects of the system would become more configurable by local admins within the target organization.	<p>The primary responsibilities are:</p> <ul style="list-style-type: none"> • Establish governance policies regarding which types of information are encouraged to be included. • Establish governance policies regarding which types of information are forbidden to be included. • Determining supported sources for data input. • Determining the privacy level and/or sharing policies available. • Monitoring usage of the system for abuse, key metrics & user satisfaction.

Goal T1: System intelligibility for decision making

This Goal applies to AI systems when the intended use of the generated outputs is to inform decision making by or about people. If this Goal applies to the system, complete the table below.

Who will use the outputs of the system to make decisions?	Who will decisions be made about?
N/A	Future capabilities may allow adjacent JobAssist AI functionality (e.g. 2A) to incorporate query results from this Knowledge Base into that workflow, thereby potentially impacting the decision-making process.

Goal T2: Communication to stakeholders

This Goal applies to all AI systems. Complete the table below.

Who will make decisions about whether to employ the system for particular tasks?	Who develops or deploys systems that integrate with this system?
Ultimately, the Job Coach will make the decision whether or not to use the system. However, it cannot be ruled out that use of the system may become mandatory per organizational policy.	The JobAssistAI team would play a role in initial integrations while a development/integration framework is created. Ultimately, anyone assigned to develop or deploy related systems on behalf of these agencies would be included here.

Goal T3: Disclosure of AI interaction

This Goal applies to AI systems that impersonate interactions with humans, unless it is obvious from the circumstances or context of use that an AI system is in use, and AI systems that generate or manipulate image, audio, or video content that could falsely appear to be authentic. If this Goal applies to the system, complete the table below.

Who will use or be exposed to the system?
Per the current scope, only trained Job Coaches will directly use the system. All AI interactions are accompanied by associated "bot" iconography and "AI" text headings.

Fairness considerations

2.4 For each Fairness Goal that applies to the system, 1) identify the relevant stakeholder(s) (e.g., system user, person impacted by the system); 2) identify any demographic groups, including marginalized groups, that may require fairness considerations; and 3) prioritize these groups for fairness consideration and explain how the fairness consideration applies. If the Fairness Goal does not apply to the system, enter "N/A" in the first column.

Goal F1: Quality of service

This Goal applies to AI systems when system users or people impacted by the system with different demographic characteristics might experience differences in quality of service that can be remedied by building the system differently. If this Goal applies to the system, complete the table below describing the appropriate stakeholders for this intended use.

Which stakeholder(s) will be affected?	For affected stakeholder(s) which demographic groups are you prioritizing for this Goal?	Explain how each demographic group might be affected.

Job Coaches	For the purposes of the initial demo release, our presumptive Job Coach is based on the information we were provided for the hackathon. We are fully aware that each Job Coach is an individual first and we must always seek to meet the user where they are.	The quality of service for Job Coaches is critical to the adoption of this feature. Over-optimizing the system for a hypothetical Job Coach will likely under-serve individual users. It will be important to revisit our assumptions based on user interviews.
-------------	--	---

Goal F2: Allocation of resources and opportunities

This Goal applies to AI systems that generate outputs that directly affect the allocation of resources or opportunities relating to finance, education, employment, healthcare, housing, insurance, or social welfare. If this Goal applies to the system, complete the table below describing the appropriate stakeholders for this intended use.

Which stakeholder(s) will be affected?	For affected stakeholder(s) which demographic groups are you prioritizing for this Goal?	Explain how each demographic group might be affected.
N/A		

Goal F3: Minimization of stereotyping, demeaning, and erasing outputs

This Goal applies to AI systems when system outputs include descriptions, depictions, or other representations of people, cultures, or society. If this Goal applies to the system, complete the table below describing the appropriate stakeholders for this intended use.

Which stakeholder(s) will be affected?	For affected stakeholder(s) which demographic groups are you prioritizing for this Goal?	Explain how each demographic group might be affected.
Job Consumers	Disabled Job Consumers, i.e. job consumers with one or more impairments relating to their vision, mobility, mental status, etc.	We consider this stakeholder group to be a vulnerable population and, accordingly, we need take precautions to minimize any stereotyping of persons with disabilities, instead focusing on the abilities of the individual job consumer. To that end, we need to consider what this process do if it encounters potentially improper language or comments in submitted material.

Technology readiness assessment

2.5 Indicate with an "X" the description that best represents the system regarding this intended use.

Select one	Technology Readiness
	The system includes AI supported by basic research and has not yet been deployed to production systems at scale for similar uses.
X	The system includes AI supported by evidence demonstrating feasibility for uses similar to this intended use in production systems.
	This is the first time that one or more system component(s) are to be validated in relevant environment(s) for the intended use. Operational conditions that can be supported have not yet been completely defined and evaluated.
	This is the first time the whole system will be validated in relevant environment(s) for the intended use. Operational conditions that can be supported will also be validated. Alternatively, nearly similar systems or nearly similar methods have been applied by other organizations with defined success.
	The whole system has been deployed for all intended uses , and operational conditions have been qualified through testing and uses in production.

Task complexity

2.6 Indicate with an "X" the description that best represents the system regarding this intended use.

Select One	Task Complexity
	Simple tasks , such as classification based on few features into a few categories with clear boundaries. For such decisions, humans could easily agree on the correct answer, and identify mistakes made by the system. For example, a natural language processing system that checks spelling in documents.
X	Moderately complex tasks , such as classification into a few categories that are subjective. Typically, ground truth is defined by most evaluators arriving at the same answer. For example, a natural language processing system that autocompletes a word or phrase as the user is typing.

	Complex tasks , such as models based on many features, not easily interpretable by humans, resulting in highly variable predictions without clear boundaries between decision criteria. For such decisions, humans would have a difficult time agreeing on the best answer, and there may be no clearly incorrect answer. For example, a natural language processing system that generates prose based on user input prompts.
--	--

Role of humans

2.7 Indicate with an "X" the description that best represents the system regarding this intended use.

Select One	Role of humans
	People will be responsible for troubleshooting triggered by system alerts but will not otherwise oversee system operation. For example, an AI system that generates keywords from unstructured text alerts the operator of errors, such as improper format of submission files.
	The system will support effective hand-off to people but will be designed to automate most use. For example, an AI system that generates keywords from unstructured text that can be configured by system admins to alert the operator when keyword generation falls below a certain confidence threshold.
X	The system will require effective hand-off to people but will be designed to automate most use. For example, an AI system that generates keywords from unstructured text alerts the operator when keyword generation falls below a certain confidence threshold (regardless of system admin configuration).
	People will evaluate system outputs and can intervene before any action is taken: the system will proceed unless the reviewer intervenes. For example, an AI system that generates keywords from unstructured text will deliver the generated keywords for operator review but will finalize the results unless the operator intervenes.
	People will make decisions based on output provided by the system: the system will not proceed unless a person approves. For example, an AI system that generates keywords from unstructured text but does not finalize the results without review and approval from the operator.

Deployment environment complexity

2.8 Indicate with an "X" the description that best represents the system regarding this intended use.

Select One	Deployment environment complexity
------------	-----------------------------------

	Simple environment , such as when the deployment environment is static, possible input options are limited, and there are few unexpected situations that the system must deal with gracefully. For example, a natural language processing system used in a controlled research environment.
X	Moderately complex environment , such as when the deployment environment varies, unexpected situations the system must deal with gracefully may occur, but when they do, there is little risk to people, and it is clear how to effectively mitigate issues. For example, a natural language processing system used in a corporate workplace where language is professional and communication norms change slowly.
	Complex environment , such as when the deployment environment is dynamic, the system will be deployed in an open and unpredictable environment or may be subject to drifts in input distributions over time. There are many possible types of inputs, and inputs may significantly vary in quality. Time and attention may be at a premium in making decisions and it can be difficult to mitigate issues. For example, a natural language processing system used on a social media platform where language and communication norms change rapidly.

Section 3: Adverse impact

Restricted Uses

3.1 *If any uses of the system are subject to a legal or internal policy restriction, list them here, and follow the requirements for those uses.*

Restricted Uses

The existing system hosting this functionality is purely hypothetical for the purposes of this demo release.

However, the presumptive customers of this system are generally part of state or local governments and will almost certainly have specific policies to which this system must adhere.

Those policies should be documented here prior to any deployment.

Unsupported uses

3.2 *Uses for which the system was not designed or evaluated or that should be avoided.*

Unsupported uses

The existing system hosting this functionality is purely hypothetical for the purposes of this demo release.

However, direct use by anyone other than Job Coaches or members of their supervisory organization that have received adequate training on the implemented features and the specific impact of these features within their environment would be unsupported in any production deployment.

Known limitations

3.3 *Describe the known limitations of the system. This could include scenarios where the system will not perform well, environmental factors to consider, or other operating factors to be aware of.*

Known limitations

The existing system hosting this functionality is purely hypothetical for the purposes of this demo release.

The system will be inherently limited by:

- The availability of, and effective access to, documents regarding policies & procedures for Job Coaches & their organization.
- The integrations established between any existing systems and JobAssistAI for the purposes of identifying user context, programmatic access to relevant data sources.
- The adoption of features by the organization and its individual users.
- Policy and/or privacy restrictions in the customer's environment

Potential impact of failure on stakeholders

3.4 Define predictable failures, including false positive and false negative results for the system as a whole and how they would impact stakeholders for each intended use.

Potential impact of failure on stakeholders

With product features directly aligned to end-user pain points, we expect rapid adoption of this system.

Therefore, any system failure will have a magnified impact. We have identified the following predictable failures:

- **System Outage:** a broad system outage (e.g. infrastructure issue) could lead to substantial reductions in efficiency as users will be trained to leverage the system in favor of any alternative processes. Interactions with, or on behalf of, Job Consumers during the outage will likely be compromised.
- **Component Outage:** the outage of one or more specific components/features could significantly impact user efficiency.
- **Data Loss:** given the acute focus on standardizing data collection, the loss of this data would have significant implications.
- **Data Theft:** given the acute focus on standardizing data collection, the theft of this data could have significant impacts on the stakeholders, especially the Job Consumers.

Potential impact of misuse on stakeholders

3.5 Define system misuse, whether intentional or unintentional, and how misuse could negatively impact each stakeholder. Identify and document whether the consequences of misuse differ for marginalized groups. When serious impacts of misuse are identified, note them in the summary of impact as a potential harm.

Potential impact of misuse on stakeholders

System misuse would be characterized as:

- Any use of the system by an unsupported or untrained user.
- The use of any system features/components in a manner contrary to or inconsistent with the training.

Sensitive Uses

3.6 Consider whether the use or misuse of the system could meet any of the Microsoft Sensitive Use triggers below.

Yes or No	Sensitive Use triggers
Yes	<p>Consequential impact on legal position or life opportunities</p> <p>The use or misuse of the AI system could affect an individual's: legal status, legal rights, access to credit, education, employment, healthcare, housing, insurance, and social welfare benefits, services, or opportunities, or the terms on which they are provided.</p>

No	Risk of physical or psychological injury The use or misuse of the AI system could result in significant physical or psychological injury to an individual.
No	Threat to human rights The use or misuse of the AI system could restrict, infringe upon, or undermine the ability to realize an individual's human rights. Because human rights are interdependent and interrelated, AI can affect nearly every internationally recognized human right.

Section 4: Data Requirements

Data requirements

4.1 Define and document data requirements with respect to the system's intended uses, stakeholders, and the geographic areas where the system will be deployed.

Data requirements

For the purposes of this demo release, only synthetic data is used in regards to individuals.

The designed use case is to support Job Coaches within the US who generally operate under a local or state government entity. Correspondingly, the system will be deployed in a US region, the specific region being chosen based on latency requirements, feature availability and/or cost considerations.

Existing data sets

4.2 If you plan to use existing data sets to train the system, assess the quantity and suitability of available data sets that will be needed by the system in relation to the data requirements defined above. If you do not plan to use predefined data sets, enter "N/A" in the response area.

Existing data sets

For the purposes of this demo release, only synthetic data is used in regards to individuals.

As the presumptive customers are state & local government agencies, access to the policies, procedures, forms, definitions, etc. is generally unrestricted.

Other datasets potential in future scope:

- State Job Placement historical data
- Job Dataset (physical/mental requirements, responsibilities, title, description, etc.)
- Business/Employer Directories

Section 5: Summary of Impact

Potential harms and preliminary mitigations

5.1 Gather the potential harms you identified earlier in the Impact Assessment in this table (check the stakeholder table, fairness considerations, adverse impact section, and any other place where you may have described potential harms). Use the mitigations prompts in the Impact Assessment Guide to understand if the Responsible AI Standard can mitigate some of the harms you identified. Discuss the harms that remain unmitigated with your team and potential reviewers.

Describe the potential harm	Corresponding Goal from the Responsible AI Standard (if applicable)	Describe your initial ideas for mitigations or explain how you might implement the corresponding Goal in the design of the system
Job Coach <i>Loss of autonomy</i> <i>Change fatigue</i>	Transparency Reliability & Safety	Comprehensive training on the features should be directly correlated with the goals of the users & their organization.
Job Consumer <i>Poor performance</i>	Reliability & Safety Fairness	Initial deployments need an enhanced degree of monitoring & auditing to assess whether the AI components are performing as intended. Effort should be taken to ensure that, in addition to remediating the system, impacted Job Consumers are made whole.
Data Loss	Reliability & Safety	The solution leverages the Azure suite of cloud services which allows a customer organization to make an informed decision regarding durability, backups and redundancy based on their risk tolerance and cost considerations.
Data Theft	Privacy & Security	The solution leverages the Azure suite of cloud services which provides world-class data privacy and security capabilities. The customer organization should seek the engagement of their Information Security partners to provide any specific requirements and perform an audit of any deployment.
Outage	Reliability & Safety	The solution leverages the Azure suite of cloud services which allows a customer organization to make an informed decision regarding failover, load-balancing, cross-region replication & many

		other reliability parameters based on their risk tolerance and cost considerations.
--	--	---

Goal Applicability

5.2 To assess which Goals apply to this system, use the tables below. When a Goal applies to only specific types of AI systems, indicate if the Goal applies to the system being evaluated in this Impact Assessment by indicating “Yes” or “No.” If you indicate that a Goal does not apply to the system, explain why in the response area. If a Goal applies to the system, you must complete the requirements associated with that Goal while developing the system.

Accountability Goals

Goals	Does this Goal apply to the system? (Yes or No)
A1: Impact assessment <i>Applies to: All AI systems.</i>	Yes
A2: Oversight of significant adverse impacts <i>Applies to: All AI systems.</i>	Yes
A3: Fit for purpose <i>Applies to: All AI systems.</i>	Yes
A4: Data governance and management <i>Applies to: All AI systems.</i>	Yes
A5: Human oversight and control <i>Applies to: All AI systems.</i>	Yes

Transparency Goals

Goals	Does this Goal apply to the system? (Yes or No)
T1: System intelligibility for decision making <i>Applies to: AI systems when the intended use of the generated outputs is to inform decision making by or about people.</i>	Yes
T2: Communication to stakeholders <i>Applies to: All AI systems.</i>	Yes

T3: Disclosure of AI interaction

Applies to: AI systems that impersonate interactions with humans, unless it is obvious from the circumstances or context of use that an AI system is in use, and AI systems that generate or manipulate image, audio, or video content that could falsely appear to be authentic.

No

If you selected "No" for any of the Transparency Goals, explain why the Goal does not apply to the system.

T3: Disclosure of AI interaction

Users will be trained on the features with specific callouts regarding the use of AI.
Any direct interaction between a User and AI will have distinctive UI elements including a "bot" profile picture for any chat and/or text headings incorporating "AI."

Fairness Goals

Goals	Does this Goal apply to the system? (Yes or No)
F1: Quality of service <i>Applies to:</i> AI systems when system users or people impacted by the system with different demographic characteristics might experience differences in quality of service that can be remedied by building the system differently.	Yes
F2: Allocation of resources and opportunities <i>Applies to:</i> AI systems that generate outputs that directly affect the allocation of resources or opportunities relating to finance, education, employment, healthcare, housing, insurance, or social welfare.	Yes
F3: Minimization of stereotyping, demeaning, and erasing outputs <i>Applies to:</i> AI systems when system outputs include descriptions, depictions, or other representations of people, cultures, or society.	Yes

If you selected "No" for any of the Fairness Goals, explain why the Goal does not apply to the system below.

Reliability & Safety Goals

Goals	Does this Goal apply to the system? (Yes or No)
RS1: Reliability and safety guidance <i>Applies to: All AI systems.</i>	Yes
RS2: Failures and remediations <i>Applies to: All AI systems.</i>	Yes
RS3: Ongoing monitoring, feedback, and evaluation <i>Applies to: All AI systems.</i>	Yes

Privacy & Security Goals

Goals	Does this Goal apply to the system? (Yes or No)
PS1: Privacy Standard compliance Applies when the Microsoft Privacy Standard applies.	Yes
PS2: Security Policy compliance Applies when the Microsoft Security Policy applies.	Yes

Inclusiveness Goal

Goals	Does this Goal apply to the system? (Yes or No)
I1: Accessibility Standards compliance Applies when the Microsoft Accessibility Standards apply.	Yes

Signing off on the Impact Assessment

5.3 Before you continue with next steps, complete the appropriate reviews and sign off on the Impact Assessment. At minimum, the PM should verify that the Impact Assessment is complete. In this case, ensure you complete the appropriate reviews and secure all approvals as required by your organization before beginning development.

Reviewer role and name	I can confirm that the document benefitted from collaborative work and different expertise within the team (e.g., engineers, designers, data scientists, etc.)	Date reviewed	Comments

Update and review the Impact Assessment at least annually, when new intended uses are added, and before advancing to a new release stage. The Impact Assessment will remain a key reference document as you work toward compliance with the remaining Goals of the Responsible AI Standard.

Scan this code to access responsible AI resources from Microsoft:



© 2022 Microsoft Corporation. All rights reserved. This document is provided "as-is." It has been edited for external release to remove internal links, references, and examples. Information and views expressed in this document may change without notice. You bear the risk of using it. Some examples are for illustration only and are fictitious. No real association is intended or inferred. This document does not provide you with any legal rights to any intellectual property in any Microsoft product. You may copy and use this document for your internal, reference purposes.